

DeepAgent: A General Reasoning Agent with Scalable Toolsets

Xiaoxi Li*
Renmin University of China
Beijing, China
xiaoxi_li@ruc.edu.cn

Yinuo Wang
Tsinghua University
Beijing, China

Ji-Rong Wen
Renmin University of China
Beijing, China
jrwen@ruc.edu.cn

Wenxiang Jiao
Jiarui Jin
Xiaohongshu Inc.
Beijing, China

Hao Wang
Xiaohongshu Inc.
Beijing, China

Yuan Lu
Xiaohongshu Inc.
Beijing, China
luyuan3@xiaohongshu.com

Guanting Dong
Jiajie Jin
Renmin University of China
Beijing, China

Yutao Zhu
Renmin University of China
Beijing, China

Zhicheng Dou†
Renmin University of China
Beijing, China
dou@ruc.edu.cn

Abstract

Large reasoning models have demonstrated strong problem-solving abilities, yet real-world tasks often require external tools and long-horizon interactions. Existing agent frameworks typically follow predefined workflows, which limit autonomous and global task completion. In this paper, we introduce **DeepAgent**, an end-to-end deep reasoning agent that performs autonomous thinking, tool discovery, and action execution within a single, coherent reasoning process. To manage long-horizon interactions, we introduce an autonomous memory folding mechanism that compresses past interactions into structured episodic, working, and tool memories, reducing error accumulation while preserving critical information. To teach general-purpose tool use efficiently and stably, we develop an end-to-end reinforcement learning strategy, namely ToolPO, that leverages LLM-simulated APIs and applies tool-call advantage attribution to assign fine-grained credit to the tool invocation tokens. Extensive experiments on eight benchmarks, including general tool-use tasks (ToolBench, API-Bank, TMDB, Spotify, ToolHop) and downstream applications (ALFWorld, WebShop, GAIA, HLE), demonstrate that DeepAgent consistently outperforms baselines across both labeled-tool and open-set tool retrieval scenarios. The code and demo are available at <https://github.com/RUC-NLPIR/DeepAgent>.

CCS Concepts

• Computing methodologies → Planning and scheduling.

Keywords

Large Reasoning Models, Autonomous Agents, Tool Retrieval, Memory Mechanism, Reinforcement Learning

*Work done during internship at Xiaohongshu Inc.

†Corresponding author.



This work is licensed under a Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License.
WWW '26, Dubai, United Arab Emirates.
© 2026 Copyright held by the owner/author(s).
ACM ISBN 979-8-4007-2307-0/2026/04
<https://doi.org/10.1145/3774904.3792460>




Figure 1: Overall performance on (a) general tool usage tasks and (b) downstream applications (best score as 100%).

ACM Reference Format:

Xiaoxi Li, Wenxiang Jiao, Jiarui Jin, Guanting Dong, Jiajie Jin, Yinuo Wang, Hao Wang, Yutao Zhu, Ji-Rong Wen, Yuan Lu, and Zhicheng Dou. 2026. DeepAgent: A General Reasoning Agent with Scalable Toolsets. In *Proceedings of the ACM Web Conference 2026 (WWW '26)*, April 13–17, 2026, Dubai, United Arab Emirates. ACM, New York, NY, USA, 12 pages. <https://doi.org/10.1145/3774904.3792460>

1 Introduction

The rapid advancement of Large Language Models (LLMs) has inspired the development of LLM-powered agents, which have found broad applications in scenarios such as web information seeking, software engineering, and personal assistance [19, 39]. Most existing agents follow predefined workflows (e.g., ReAct [68] and Plan-and-Solve [52]) with iterative “Reason-Act-Observe” loops (Figure 2(a)). Although effective in simpler tasks, these approaches suffer from several critical limitations: (1) lack of autonomy in execution steps and overall procedure; (2) inability to dynamically discover tools during task execution; (3) deficiency in fully autonomous management of interactive memory; and (4) insufficient depth and coherence in reasoning about the entire task. These limitations hinder agents from real-world problems, particularly for complex tasks that demand general and multiple tool-use.




Figure 2: Comparison of agent paradigms: (a) Traditional agents with predefined workflows, (b) Deep Research agents that can autonomously call limited tools, and (c) Our DeepAgent, a fully autonomous reasoning agent that dynamically discovers and invokes helpful tools, all within a continuous agentic reasoning process.

Recently, the advent of Large Reasoning Models (LRMs) has demonstrated the capability to solve complex problems in domains like mathematics, programming, and scientific reasoning through a step-by-step “slow thinking” process [2, 30, 55]. However, many real-world tasks necessitate the use of external tools for their completion. Recent approaches integrate tool use into reasoning [25, 29, 72], but typically rely on a small, fixed tool set such as search, browsing, and coding (Figure 2(b)), limiting their generality.

To address these challenges, we introduce **DeepAgent**, an end-to-end deep reasoning agent that can complete an entire task by dynamically retrieving and calling tools within a single, coherent agentic reasoning process. As depicted in Figure 2(c), DeepAgent operates by autonomously thinking, searching for tools, and executing actions. This paradigm shifts away from traditional, predefined workflows that rely on predefined tools, task planning, and iterative tool use. Instead, DeepAgent maintains a global perspective on the entire task, unconstrained by the need to deliberate on specific, isolated operations. Tools are not pre-retrieved in advance but are dynamically discovered on an as-needed basis, thereby fully unlocking the autonomous potential of the large reasoning model.

To facilitate robust exploration in long-horizon environments, we equip DeepAgent with **Autonomous Memory Folding**. This strategy allows the agent to dynamically consolidate its reasoning process and interaction history into a *structured memory schema*. Beyond reducing token overhead, this mechanism enables the agent to “take a breath”—pausing to reconsider strategies and avoid erroneous paths. To minimize information loss during consolidation, we introduce a *brain-inspired memory architecture* comprising episodic, working, and tool memory, all structured with an agent-usable data schema to ensure the stability and utility of the folded memory.

To enhance DeepAgent’s proficiency in mastering these mechanisms, we propose **ToolIPO**, an end-to-end reinforcement learning (RL) training method tailored for general tool use. Existing agentic RL training in general domains presents two significant challenges: (1) The reliance on a multitude of real-world APIs during training can lead to instability, slow execution, and high costs. To prevent this, we leverage *LLM-simulated APIs*, which enhance the stability and efficiency of the training process. (2) A sparse reward based solely on the final outcome is often insufficient to guarantee the

accuracy of intermediate tool calls. We address this by implementing *tool-call advantage attribution*, which precisely assigns credit to the specific tokens responsible for correct tool invocations, thereby providing a more granular and effective learning signal.

We conduct extensive experiments on a wide range of benchmarks. For **(1) General Tool-Use Tasks**, we evaluate DeepAgent on ToolBench, API-Bank, TMDB, Spotify, and ToolHop, which feature toolsets scaling from tens to over ten thousand distinct tools. For **(2) Downstream Applications**, we test its performance on ALFWorld, WebShop, GAIA, and Humanity’s Last Exam (HLE), which require the use of domain-specific toolsets. The overall results in Figure 1 show that DeepAgent achieves superior performance across all scenarios.

Our main contributions are summarized as follows:

- (1) We propose DeepAgent, the first agentic framework that enables reasoning models to autonomously think, discover tools, and execute actions within a unified reasoning process, empowering LRMs to harness toolsets of arbitrary scale and generalize to complex real-world tasks.
- (2) We introduce an autonomous memory folding mechanism, complemented by a brain-inspired memory design. This endows the agent with the ability to “take a breath” and reconsider its exploration strategies following unsuccessful attempts.
- (3) We propose an end-to-end reinforcement learning training methodology for general-purpose tool use, ensuring stability and efficiency in large-scale tool execution during training, as well as accuracy in tool invocation during reasoning.
- (4) We conduct extensive experiments across eight benchmarks, demonstrating DeepAgent’s superior tool-use capabilities and high adaptability to real-world tasks.

2 Related Work

2.1 Large Reasoning Models

Large Reasoning Models (LRMs) [4, 16] have demonstrated significant performance improvements in mathematical, scientific, and coding tasks by employing step-by-step slow thinking processes before generating final responses. Existing research has explored




Figure 3: Overview of the DeepAgent framework. The main reasoning model autonomously discovers tools, executes actions, and folds previous memory to restart with structured memories, all within a unified thinking process. The DeepAgent is trained end-to-end with ToolPO, an RL method that uses a tool simulator to simulate large-scale real-world tool APIs, and rewards both final task success and correct intermediate tool calls through fine-grained advantage attribution.

various approaches to elicit extended Chain-of-Thought (CoT) reasoning [60] from models, including data synthesis for Supervised Fine-Tuning (SFT) [36, 54], and end-to-end RL [4]. Additionally, substantial work has investigated optimization strategies for reasoning models, such as advanced RL training algorithms [58] and improving reasoning efficiency [66]. However, models relying solely on parametric knowledge face inherent limitations and cannot interact with the real world. Recent studies have begun exploring tool-augmented reasoning approaches, including Search-o1 [25], Search-R1 [18], ToRL [29], DeepResearcher [72], and SimpleTIR [64]. However, these methods typically support only a limited set of research-oriented tools, such as web search, page browsing, and code execution, which constrains their applicability to real-world scenarios that demand access to more diverse tools.

2.2 Autonomous Agents

LLM-powered autonomous agents accomplish real-world tasks by invoking external tools to interact with their environment [6, 7, 13–15, 20, 22, 27, 28, 31, 38, 46, 49, 57, 61, 71]. Current agent methodologies, including ReAct [68], Plan-and-Solve [52], Reflexion [44], and CodeAct [56], predominantly follow predefined workflows with fixed execution patterns. This rigid structure limits their ability to fully leverage the autonomous decision-making and deep reasoning capabilities of advanced reasoning models. Recent efforts have investigated training LLMs to autonomously invoke tools through data synthesis and SFT methods [9, 63] and RL training frameworks [3, 5, 8, 10, 17, 23, 32, 48, 59]. However, most existing methods rely on pre-selected, labeled tools, which limit their applicability to real-world scenarios. Real-world tasks are highly variable and require access to diverse toolsets that cannot be predetermined, aligning with the emerging Model Context Protocol (MCP) [12] paradigm. Although some prior work has explored tool retrieval

mechanisms [37, 42, 53], most approaches conduct only a single upfront retrieval step and incorporate the retrieved tools, with limited exploration of dynamic tool discovery during task execution. Therefore, we aim to develop a deep reasoning agent capable of dynamically discovering and invoking helpful tools from scalable toolsets to address more generalized real-world tasks.

3 Methodology

3.1 Problem Formulation

We frame the agent's task as a sequential decision-making process. The agent receives a user-provided question Q and an instruction I , and interacts with an environment over a series of steps $t = 1, \dots, T$ to accomplish the specified goal. The environment provides access to a collection of tools $\mathcal{T}$ at an arbitrary scale.

At each step t , the agent's state s_t consists of the history of all previous actions and their resulting observations, i.e., $s_t = (a_1, o_1, \dots, a_{t-1}, o_{t-1})$. The agent, driven by a policy π parameterized by θ , selects an action a_t based on the current state, the user question, and the instruction:

$$a_t \sim \pi_\theta(\cdot | s_t, Q, I). \quad (1)$$

An action a_t can be one of four types:

- **Internal Thought (a_t^{think}):** A textual reasoning step generated by the LRM to analyze the problem or plan its next steps. The corresponding observation o_t is typically empty.
- **Tool Search (a_t^{search}):** A natural language query q_s to find relevant tools from $\mathcal{T}$. The observation o_t is a list of retrieved tools.
- **Tool Call (a_t^{call}):** The invocation of a specific tool $\tau \in \mathcal{T}$ with a set of arguments. The observation o_t is the execution result returned by the tool.

- **Memory Fold** (a_t^{fold}): A special action to compress the interaction history s_t into a structured memory summary. The subsequent state s_{t+1} is then initialized with this compressed memory. The sequence of states, actions, and observations forms a trajectory $\tau = (s_1, a_1, o_1, \dots, s_T, a_T, o_T)$. The process terminates when the agent completes the task or reaches a maximum step limit. Suppose $R(\tau)$ is a reward function that evaluates the overall success of the trajectory τ , the objective is to learn an optimal policy π_θ^* that maximizes the expected cumulative reward for a given task:

$$\pi_\theta^* = \arg \max_{\pi_\theta} \mathbb{E}_{\tau \sim \pi_\theta} [R(\tau)]. \quad (2)$$

3.2 Overview of the DeepAgent Framework

As illustrated in Figure 3, the DeepAgent framework is architected around a main reasoning process, which is supported by several auxiliary mechanisms to ensure robustness and efficiency.

- **Main Reasoning Process:** The core of DeepAgent is a powerful large reasoning model that drives the entire task-completion process. In a single stream of thought, the LRM autonomously reasons about the task, dynamically discovers necessary tools, executes actions, and manages its own memory. This unified approach departs from traditional, rigid agent workflows, allowing the LRM to maintain a global perspective on the task.
- **Auxiliary Mechanisms:** DeepAgent employs an auxiliary LLM to handle complex interactions with large toolsets and manage long histories. This background model enhances system stability by: (1) filtering and summarizing retrieved tool documentation if it's too lengthy, (2) denoising and condensing verbose information returned from tool calls, and (3) compressing long interaction histories into a structured memory. This division of labor allows the main LRM to concentrate on high-level strategic reasoning.

3.3 Autonomous Tool Search and Calling

DeepAgent's main LRM performs all actions by generating specific textual prompts within its continuous reasoning process. These actions are then intercepted and executed by the system.

Tool Search. When the agent determines it needs a tool, it generates a tool search query q_s encapsulated within special tokens: `<tool_search>  $q_s$  </tool_search>`. The system's tool retriever operates via dense retrieval. First, we build an index by pre-computing an embedding $E(d_i)$ for the documentation d_i of each tool $\tau_i \in \mathcal{T}$ using an embedding model E . During inference, given the query q_s , the system retrieves the top- k tools by ranking them based on the cosine similarity $\text{sim}(\cdot, \cdot)$:

$$\mathcal{T}_{\text{retrieved}} = \underset{\tau_i \in \mathcal{T}}{\text{top-k}} (\text{sim}(E(q_s), E(d_i))). \quad (3)$$

The retrieved tool documentation is then processed by the auxiliary LLM—summarized if too lengthy, otherwise provided directly—and returned to the main LRM's context: `<tool_search_result> relevant tools </tool_search_result>`.

Tool Call. To execute a tool, the agent generates a structured call including the tool's name and arguments: `<tool_call> {"name": "tool_name", "arguments": ...} </tool_call>`. The framework parses this call, executes the tool, and captures the output. This output is,

if necessary, summarized by the auxiliary LLM to ensure it is concise and helpful, before being fed back into the reasoning context: `<tool_call_result> helpful information </tool_call_result>`.

3.4 Autonomous Memory Folding and Brain-Inspired Memory Schema

The agent can trigger memory folding at any logical point in its reasoning process—such as after completing a sub-task or realizing an exploration path was incorrect—by generating a special token: `<fold_thought>`. Upon detecting this token, the system initiates the memory folding process. The auxiliary LLM (parameterized by θ_{aux}) processes the entire preceding interaction history s_t and generates three structured memory components in parallel:

$$(M_E, M_W, M_T) = f_{\text{compress}}(s_t; \theta_{\text{aux}}). \quad (4)$$

These compressed episodic (M_E), working (M_W), and tool (M_T) memories then replace the raw interaction history, enabling the agent to proceed with a refreshed and condensed view of its progress while avoiding entrapment in incorrect exploration paths.

Inspired by human cognitive systems, the structured memory M_t is composed of three distinct components that are generated in parallel: $M_t = (M_E, M_W, M_T)$, where M_E, M_W, M_T denote episodic, working, and tool memories, respectively.

- **Episodic Memory (M_E):** This component serves as a high-level log of the task, recording key events, major decision points, and sub-task completions. It provides the agent with long-term context regarding the overall task structure and its overarching goals.
- **Working Memory (M_W):** This contains the most recent information, such as the current sub-goal, obstacles encountered, and near-term plans. It is the core component that ensures the continuity of the agent's reasoning across the memory fold.
- **Tool Memory (M_T):** This consolidates all tool-related interactions, including which tools have been used, how they were invoked, and their effectiveness. It allows the agent to learn from its experiences, refining its tool selection and usage strategies.

To ensure that the compressed memory is stable and easily parsed by the agent, we employ an **agent-usable data schema** in JSON format instead of unstructured natural language. It offers two main benefits: maintaining a controllable and predictable structure, and mitigating the loss of critical details that can occur when summarizing long-form text. Details of the data schema are in Appendix C.

3.5 End-to-end RL Training with ToolPO

We train DeepAgent end-to-end with Tool Policy Optimization (ToolPO), an RL approach designed for general tool-using agents.

Training Data Collection. We first collect a diverse training dataset spanning four categories. To instill **general tool-use** capabilities, we use ToolBench [37]. For **real-world interaction**, we leverage ALFWorld [45] and WebShop [67]. To enhance **deep research** skills, we incorporate data from WebDancer [61] and WebShaperQA [50]. Lastly, to improve **mathematical reasoning** with code, we use DeepMath [11]. Further details are available in Appendix A.1.

Tool Simulator. Training an agent that interacts with thousands of real-world APIs is often impractical due to instability, latency, and cost. To address this, we develop an **LLM-based Tool Simulator**.

Table 1: Main results on general tool usage tasks, encompassing scenarios with both labeled tools and open-set tool retrieval over large-scale toolsets. We report Pass@1 metric for all tasks. For 32B models, the best results are in bold and the second are underlined. Results from larger or closed-sourced models are in gray color for reference.

Method	Backbone	ToolBench		API-Bank		TMDB		Spotify		ToolHop	
		Success	Path	Success	Path	Success	Path	Success	Path	Correct	Path
Scenario 1: Completing Tasks w/ Ground-truth Tools											
Workflow-based Methods											
ReAct	Qwen2.5-32B	41.0	64.7	60.4	68.3	46.0	65.3	29.8	56.3	37.6	49.1
CodeAct	Qwen2.5-32B	53.0	68.3	62.4	70.6	48.0	67.4	33.3	58.7	34.7	48.8
Plan-and-Solve	Qwen2.5-32B	52.0	65.4	58.4	67.5	51.0	71.6	28.1	54.8	39.2	49.7
ReAct	QwQ-32B	52.0	61.6	73.3	78.6	43.0	65.3	47.4	69.4	47.4	51.6
CodeAct	QwQ-32B	54.0	63.4	74.3	79.4	55.0	74.5	52.6	75.4	43.2	53.4
Plan-and-Solve	QwQ-32B	55.0	64.7	70.3	75.4	48.0	61.3	49.1	70.6	45.4	50.6
ReAct	Qwen2.5-72B	56.0	69.3	73.3	78.6	47.0	67.7	57.9	76.6	44.8	55.4
ReAct	GPT-4o	52.0	53.9	79.2	83.3	77.0	89.3	47.4	70.6	40.0	53.7
ReAct	DeepSeek-R1	57.0	68.3	71.3	76.2	76.0	89.0	64.9	81.3	50.2	61.8
Autonomous Tool Usage within Reasoning											
DeepAgent-32B-Base	QwQ-32B	63.0	74.3	76.2	81.0	85.0	92.0	70.2	89.3	49.1	59.8
DeepAgent-32B-RL	QwQ-32B	69.0	78.6	75.3	80.2	89.0	94.8	75.4	92.0	51.3	62.5
Scenario 2: Completing Tasks w/ Open-Set Tool Retrieval											
Workflow-based Methods											
ReAct	Qwen2.5-32B	55.0	20.8	16.0	42.0	11.0	34.5	7.0	25.4	13.2	17.9
CodeAct	Qwen2.5-32B	51.0	19.0	22.0	49.6	19.0	46.8	10.5	31.6	12.7	17.4
Plan-and-Solve	Qwen2.5-32B	54.0	20.4	18.0	42.8	15.0	40.5	8.8	26.3	12.0	16.3
ReAct	QwQ-32B	44.0	19.0	20.0	52.7	18.0	40.3	22.8	45.5	27.1	22.3
CodeAct	QwQ-32B	48.0	21.6	16.0	45.0	31.0	52.8	24.6	49.6	29.0	26.1
Plan-and-Solve	QwQ-32B	45.0	19.6	18.0	44.3	24.0	46.8	19.3	42.7	25.7	20.8
ReAct	Qwen2.5-72B	52.0	21.6	14.0	38.9	28.0	50.7	21.1	48.5	21.1	19.9
ReAct	GPT-4o	41.0	28.9	18.0	42.8	35.0	56.8	17.5	26.3	24.1	28.6
ReAct	DeepSeek-R1	47.0	22.3	12.0	57.3	34.0	53.1	29.8	51.7	36.2	32.9
Autonomous Tool Retrieval and Usage within Reasoning											
DeepAgent-32B-Base	QwQ-32B	60.0	35.7	22.0	61.8	52.0	71.8	49.1	68.6	38.4	40.3
DeepAgent-32B-RL	QwQ-32B	64.0	37.2	24.0	64.9	55.0	74.3	50.9	74.4	40.6	40.5

This simulator, powered by an auxiliary LLM, mimics the responses of real-world APIs (e.g., RapidAPI). This approach provides a stable, efficient, and low-cost environment for robust RL training.

Global and Tool-Call Advantage Attribution. For each input prompt, we sample a group of K trajectories $\{\tau_1, \dots, \tau_K\}$. ToolPO defines two distinct reward components. The first is a reward for overall task success, $R_{\text{succ}}(\tau)$, which is a task-success score that reflects the quality of the final outcome (e.g., the accuracy of the final answer). The second is a tool-call reward, $R_{\text{action}}(\tau)$, which reflects the quality of intermediate actions. This action-level reward is composed of rewards for correct tool invocations and efficient memory folding. Specifically, $R_{\text{action}}(\tau) = \lambda_1 \sum_{t=1}^T C(a_t^{\text{call}}) + \lambda_2 S_{\text{pref}}(\tau)$, where $C(a_t^{\text{call}})$ is 1 if a tool call is correct and 0 otherwise. $S_{\text{pref}}(\tau)$ is a preference score encouraging efficient use of memory folding, defined by comparing a trajectory with folding (τ_{fold}) to one without (τ_{direct}): $S_{\text{pref}} = (L(\tau_{\text{direct}}) - L(\tau_{\text{fold}})) / (L(\tau_{\text{direct}}) + L(\tau_{\text{fold}}))$.

Based on these rewards, we compute two separate group-relative advantages. The task success advantage for trajectory τ_k is:

$$A_{\text{succ}}(\tau_k) = R_{\text{succ}}(\tau_k) - \frac{1}{K} \sum_{j=1}^K R_{\text{succ}}(\tau_j), \quad (5)$$

which is attributed to all generated tokens in τ_k , providing a global learning signal. Similarly, the action-level advantage is:

$$A_{\text{action}}(\tau_k) = R_{\text{action}}(\tau_k) - \frac{1}{K} \sum_{j=1}^K R_{\text{action}}(\tau_j). \quad (6)$$

Crucially, this advantage is attributed *only* to the specific tokens that constitute the tool call and memory folding actions. This fine-grained credit assignment provides a more targeted signal for learning correct and efficient tool use.

Optimization Objective. The total advantage for a given token y_i in trajectory τ_k is the sum of the global and local advantages:

$$A(y_i) = A_{\text{succ}}(\tau_k) + M(y_i) \cdot A_{\text{action}}(\tau_k), \quad (7)$$

where $M(y_i)$ is a mask that is 1 if y_i is part of a tool-call or memory-fold token sequence, and 0 otherwise. ToolPO then optimizes the policy using a clipped surrogate objective function:

$$\mathcal{L}_{\text{ToolPO}}(\theta) = \mathbb{E}_{\tau_k} \left[\sum_{i=1}^{|\tau_k|} \min \left(\rho_i(\theta) A(y_i), \text{clip}(\rho_i(\theta), 1 - \epsilon, 1 + \epsilon) A(y_i) \right) \right]. \quad (8)$$

Table 2: Main results on downstream task applications, spanning Embodied AI (ALFWorld), Online Shopping (WebShop), General AI Assistants (GAIA), and Humanity’s Last Exam (HLE). We report Pass@1 for all tasks. For 32B models, the best results are in bold and the second are underlined. Results from larger or closed-sourced models are in gray color for reference.

Method	Backbone	ALFWorld		WebShop		GAIA				HLE		
		Success	Path	Success	Score	Text	MM	File	All	Text	MM	All
Completing Tasks w/ Task-specific Toolsets												
Workflow-based Methods												
ReAct	Qwen2.5-32B	60.4	79.1	6.0	28.8	25.2	16.7	13.2	21.2	6.5	7.1	6.6
CodeAct	Qwen2.5-32B	65.7	83.3	12.4	34.5	28.2	20.8	18.4	24.8	7.5	8.0	7.6
Reflexion	Qwen2.5-32B	66.4	86.0	9.2	31.6	29.1	20.8	18.4	25.5	5.9	5.3	5.8
Plan-and-Solve	Qwen2.5-32B	63.4	80.4	7.6	29.3	27.2	16.7	15.8	23.0	7.2	6.2	7.0
ReAct	QwQ-32B	82.1	87.8	17.2	45.3	35.0	8.3	36.8	31.5	13.2	8.8	12.2
CodeAct	QwQ-32B	78.4	86.2	18.0	46.4	38.8	20.8	31.6	34.5	14.2	8.0	12.8
Reflexion	QwQ-32B	85.1	88.4	21.6	50.4	37.9	20.8	36.8	35.2	11.9	7.1	10.8
Plan-and-Solve	QwQ-32B	79.1	84.7	16.0	43.8	36.9	16.7	34.2	33.3	12.9	9.7	12.2
AgentLM*	Llama2-70B	86.0	-	-	64.9	-	-	-	-	-	-	-
ReAct	Qwen2.5-72B	86.5	86.5	22.0	44.5	32.0	20.8	31.6	30.3	9.0	8.0	8.8
ReAct	DeepSeek-R1	79.1	85.8	19.6	49.7	43.7	29.2	39.5	40.6	14.2	8.8	13.0
ReAct	GPT-4o	65.7	87.8	15.6	52.5	35.0	16.7	36.8	32.7	13.2	10.6	12.6
ReAct	Claude-4	93.3	91.5	20.4	56.6	56.3	37.5	52.6	52.7	15.5	16.8	15.8
Autonomous Tool Usage within Reasoning												
Deep Research	OpenAI (o3)	-	-	-	-	-	-	-	67.4	-	-	26.6
WebThinker	QwQ-32B	-	-	-	-	48.5	25.0	13.2	37.0	14.2	8.8	13.0
HiRA	QwQ-32B	84.3	87.6	23.2	51.9	44.7	33.3	42.1	42.5	14.5	10.6	13.6
DeepAgent-32B-Base	QwQ-32B	88.1	91.4	32.0	55.4	49.5	37.5	44.7	46.7	19.1	13.3	17.8
DeepAgent-32B-RL	QwQ-32B	91.8	92.0	34.4	56.3	58.3	33.3	52.6	53.3	21.7	15.0	20.2

Here, $\rho_i(\theta) = \frac{\pi_\theta(y_i|y_{<i},s)}{\pi_{\theta_{\text{old}}}(y_i|y_{<i},s)}$ is the probability ratio for token y_i . This objective encourages the model to increase the probability of both intermediate actions and end-to-end task accomplishment that exhibit positive relative advantage, thereby ensuring stable and effective policy updates.

4 Experimental Settings

4.1 Tasks and Datasets

We conduct extensive experiments on a wide range of benchmarks, including general tool-use and downstream applications.

General Tool-Use. These benchmarks cover toolsets from tens to > 10k tools, and thus stress scalability. They evaluate core capabilities for general tool use, including tool planning, tool retrieval, and accurate multi-step tool calling. We use **ToolBench** [37] (16k+ real-world APIs; G3 subset with multi-step/multi-tool calls), **API-Bank** [24] (314 dialogues; 73 APIs; 753 calls), **RestBench** [47] (TMDB: 54 tools, 2.3 calls/question; Spotify: 40 tools, 2.6 calls/question), and **ToolHop** [69] (3,912 executable tools; 3–7 calls/task). We evaluate two settings: provided ground-truth tools and open-set tool retrieval from the full toolset.

Downstream Applications. We also evaluate downstream applications with domain-specific toolsets: **ALFWorld** [45] (text embodied tasks with nine actions, e.g., move/take), **WebShop** [67] (shopping with ‘search’ and ‘click’), **GAIA** [33] (web search/browsing, VQA, code, file reading), and **Humanity’s Last Exam (HLE)** [35] (code, search, browsing, VQA). These tasks test long-horizon interaction

in more realistic environments, requiring state tracking, error recovery, and coordination across heterogeneous tools; we equip agents with task-specific toolsets.

4.2 Baselines

Our baselines include: (1) **Workflow-based Methods**: ReAct [68] alternates explicit reasoning with environment actions in a Reason-Act-Observe loop. CodeAct [56] expresses actions as executable Python code that runs in an interpreter. Plan-and-Solve [52] first sketches a high-level plan and then executes it step by step. Reflexion [43] enhances learning through verbal self-reflection after failed attempts. AgentLM [70] uses instruction tuning to enhance general agent capabilities of LLMs. (2) **Autonomous Tool Usage within Reasoning**: WebThinker [26] interleaves thinking with web search and deep web exploration. HiRA [21] introduces a hierarchical agent architecture where a meta planner decomposes tasks, a coordinator routes subtasks, and specialized executors solve them with dual-channel memory. OpenAI Deep Research [34] is an agentic system based on reasoning models.

4.3 Implementation Details

We use QwQ-32B [51] as DeepAgent’s backbone model, with Qwen2.5-32B-Instruct [40] as the auxiliary model in our main results. Text generation employs a maximum of 81,920 tokens with temperature 0.7, top_p 0.8, top_k 20, and repetition penalty 1.05. Web search and page browsing are implemented using Google Serper API and




Figure 4: Visualization of training dynamics, including (a) reward scores and (b) validation scores across training steps.

Table 3: Ablation studies on the components of DeepAgent, where the best results are in bold.

Method	Tool-Usage		Application		Avg.
	ToolB.	ToolH.	WebS.	GAIA	
DeepAgent-32B-RL	64.0	40.6	34.4	53.3	48.1
w/o Training (Base)	60.0	38.4	32.0	46.7	44.3
w/o Memory Folding	63.0	36.6	32.4	44.7	44.2
w/o Tool Simulation	62.0	35.2	33.6	48.5	44.8
w/o Tool Adv. Attribution	62.0	39.6	33.2	49.5	46.1

Jina Reader API, respectively. The VQA tool is based on Qwen2.5-VL-32B-Instruct [1]. Tool retrieval is performed using bge-large-en-v1.5 [62]. Training consists of 100 steps of ToolPO with batch size 64, $\lambda_1 = \lambda_2 = 1$, rollout size $K = 8$, and maximum sequence length 32,768. Additional details are provided in Appendix B. All experiments are conducted on 64 NVIDIA H20-141GB GPUs.

5 Experimental Results

5.1 Main Results on General Tool Usage Tasks

Table 1 summarizes results on general tool-use tasks and yields three observations. (1) **DeepAgent’s End-to-End Reasoning Surpasses Workflow-Based Methods.** DeepAgent consistently outperforms workflow-based agents. On labeled-tool tasks, DeepAgent-32B-RL reaches 89.0% on TMDb and 75.4% on Spotify, exceeding the best 32B baselines (55.0% and 52.6%). This highlights the advantage of end-to-end agentic reasoning over rigid, predefined action loops. (2) **DeepAgent Maintains Robustness in Open-Set Scenarios.** Gains are larger in open-set settings where tool discovery is required: on ToolBench and ToolHop, DeepAgent-32B-RL achieves 64.0% and 40.6%, surpassing the best baselines (54.0% and 29.0%). This suggests that on-demand tool discovery within the reasoning process is both more robust and more scalable in realistic open-set tool environments. (3) **ToolPO Training Further Improves Tool-Usage Capabilities.** ToolPO yields consistent improvements over the base model, increasing ToolBench success by up to 6.0% and Spotify (labeled) by 5.2%. These gains indicate that our RL training better aligns intermediate tool calls with end-task success.

Table 4: Effectiveness analysis of autonomous tool retrieval strategy in open-set scenarios compared to pre-retrieved tool methods. Numbers in parentheses indicate toolset sizes.

Method	ToolB. (16k)	ToolH. (3.9k)	TMDb (54)	Spotify (40)	Avg.
ReAct Workflow					
Input Retrieved Tool	35.0	25.4	14.0	15.0	22.4
Auto. Tool Retrieval	34.0	37.1	18.0	27.8	28.0
Plan-and-Solve Workflow					
Input Retrieved Tool	37.0	24.8	19.0	16.0	24.2
Auto. Tool Retrieval	45.0	25.7	24.0	19.3	28.5
End-to-end Agentic Reasoning (DeepAgent)					
Input Retrieved Tool	53.0	37.0	34.0	43.9	42.0
Auto. Tool Retrieval	64.0	40.6	55.0	50.9	52.6

5.2 Main Results on Downstream Applications

Table 2 reports the downstream results that require long-horizon interaction and more complex environment dynamics. (1) **The autonomous reasoning paradigm generally outperforms the workflow-based methods.** Methods that integrate tool use into continuous reasoning outperform workflow-based agents. On GAIA, DeepAgent-32B-Base (46.7) and HiRA (42.5) exceed the best workflow baseline CodeAct (34.5); on WebShop, DeepAgent-32B-Base scores 32.0 vs. 18.0. This supports that long-horizon tasks benefit from flexible, integrated reasoning-and-action rather than fixed workflows. (2) **DeepAgent demonstrates superior performance across various application tasks.** DeepAgent achieves the best performance among 32B models: 53.3 on GAIA (vs. 42.5 for HiRA) and 91.8% on ALFWorld (vs. 84.3). We attribute this to DeepAgent’s coherent reasoning process and its support for robust long-horizon interaction (e.g., autonomous memory folding). (3) **ToolPO training further improves performance on downstream applications.** ToolPO further improves downstream performance: GAIA 46.7 $\rightarrow$ 53.3 (+6.6) and ALFWorld 88.1% $\rightarrow$ 91.8% (+3.7). This shows the tool-use improvements learned by ToolPO transfer to interactive downstream settings.

5.3 Analysis of Training Dynamics

Figure 4 shows the training dynamics of DeepAgent, including the reward scores and validation scores across training steps. As shown in the figure, (1) **DeepAgent trained with ToolPO achieves higher upper bounds on both reward and validation scores compared to the commonly used GRPO.** (2) **Moreover, the training reward exhibits less fluctuation than GRPO, demonstrating better training stability.** This indicates that using tool simulators instead of directly training with unstable real-world APIs, along with employing tool-call process supervision, enables more stable and effective training of tool-usage capabilities.

5.4 Ablation Studies

We conduct ablation studies in Table 3 to validate the effectiveness of each component in DeepAgent. (1) **Importance of ToolPO Training:** Removing ToolPO training (the Base model) results in the most significant performance drop (from 48.1 to 44.3). This




Figure 5: Scaling analysis of performance with respect to maximum action limits on WebShop and GAIA datasets.

highlights the central role of our end-to-end RL method in enhancing tool use and complex task completion. **(2) Effectiveness of Memory Folding:** The absence of memory folding also leads to a substantial performance decline (average score drops to 44.2), particularly on the long-horizon task GAIA (from 53.3 to 44.7). This confirms that the autonomous memory folding mechanism, allowing the agent to "take a breath" and replan, is crucial for robust long-term interaction. **(3) Contribution of Training Strategies:** Removing the tool simulator and tool-call advantage attribution both lead to performance degradation. This validates that the tool simulator enables more stable training, and fine-grained advantage attribution provides precise learning signals.

5.5 Effectiveness of Tool Retrieval Strategies

To compare pre-retrieving tools versus autonomous discovery during task execution, we conduct experiments shown in Table 4. We find: **(1) The on-demand nature of dynamic tool discovery yields superior performance and robust scalability.** Autonomous tool retrieval during reasoning consistently outperforms pre-retrieved tools across all frameworks, demonstrating the superiority of on-demand tool access in open-set scenarios. Performance gains are most pronounced on large toolsets like ToolBench (16k tools) and ToolHop (3.9k tools), indicating robust scalability for real-world tasks. **(2) DeepAgent synergizes better with dynamic retrieval.** Combined with autonomous tool retrieval, our framework achieves the best results by a large margin, scoring 52.6 on average versus 28.5 for the best workflow-based method. This demonstrates that DeepAgent’s architecture is uniquely suited for dynamic tool discovery.

5.6 Scaling Analysis of Action Limits

Figure 5 illustrates the performance of DeepAgent and ReAct on the WebShop and GAIA datasets as the maximum action limit is varied. The results yield several key insights. **(1) DeepAgent consistently and significantly outperforms the ReAct baseline across all tested action limits on both datasets,** demonstrating its superior effectiveness. **(2) For both agents, performance generally improves as the maximum number of actions increases.** This suggests that complex tasks benefit from a longer interaction horizon, allowing for more thorough exploration and reasoning. **(3) DeepAgent exhibits stronger scalability.** As the action limit increases, the performance gap between DeepAgent

Table 5: Performance with different reasoning model backbones: MOE-based models with 30B and 235B parameters.

Method	Tool-Usage		Application			Avg.
	ToolB.	ToolH.	ALF.	WebS.	GAIA	
<i>Qwen3-30B-A3B-Thinking</i>						
ReAct	52.0	22.0	67.9	18.4	34.5	35.7
Plan-and-Solve	50.0	23.6	68.7	20.4	35.2	37.0
DeepAgent (Base)	59.0	47.5	69.4	31.4	39.4	46.9
<i>Qwen3-235B-A22B-Thinking</i>						
ReAct	61.0	40.9	79.9	21.6	36.4	45.1
Plan-and-Solve	63.0	43.0	78.4	24.4	38.4	46.0
DeepAgent (Base)	67.0	48.2	85.8	37.2	51.5	55.7

and ReAct widens, particularly on WebShop. This sustained gain suggests DeepAgent strategically selects effective, task-relevant actions, avoiding the wasteful steps that limit ReAct’s scalability.

5.7 Generalization Across Different Backbones

Table 5 shows the performance of DeepAgent with different backbone large reasoning models, including Qwen3-30B-A3B-Thinking and Qwen3-235B-A22B-Thinking [65]. **(1) DeepAgent consistently outperforms workflow-based methods.** With both the 30B and 235B MoE-based reasoning models as backbones, DeepAgent maintains a significant performance margin over ReAct and Plan-and-Solve, demonstrating the generalizability of its agentic reasoning approach. **(2) DeepAgent scales effectively with larger models.** While all methods benefit from scaling the backbone from a 30B to a 235B model, DeepAgent shows the largest absolute performance gains on complex application tasks.

6 Conclusion

In this work, we introduce DeepAgent, an end-to-end reasoning agent that unifies thinking, tool discovery, and execution into a single, coherent agentic reasoning process. To enable robust long-horizon interaction, we propose an autonomous memory folding mechanism that compresses interaction history into a structured memory, allowing the agent to "take a breath" and reconsider its strategy. We also introduce ToolPO, an end-to-end RL method that leverages LLM simulated APIs for stable training and fine-grained advantage attribution for precise credit assignment to tool invocations. Extensive experiments on general tool-use and downstream applications demonstrate that DeepAgent significantly outperforms various baseline agents, particularly in open-set scenarios requiring dynamic tool discovery over scalable toolsets. This work opens new avenues for developing more general and scalable LLM agents for broader real-world applications.

Acknowledgments

This work was supported by the National Natural Science Foundation of China No. 62272467, and the China Postdoctoral Science Foundation under Grant Number 2025T180440. The work was partially done at the Engineering Research Center of Next-Generation Intelligent Search and Recommendation, MOE.

References

- [1] Shuai Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, Sibao Song, Kai Dang, Peng Wang, Shijie Wang, Jun Tang, Humen Zhong, Yuanzhi Zhu, Ming-Hsuan Yang, Zhaohai Li, Jianqiang Wan, Pengfei Wang, Wei Ding, Zheren Fu, Yiheng Xu, Jiabo Ye, Xi Zhang, Tianbao Xie, Zesen Cheng, Hang Zhang, Zhibo Yang, Haiyang Xu, and Junyang Lin. 2025. Qwen2.5-VL Technical Report. *CoRR* abs/2502.13923 (2025). arXiv:2502.13923 doi:10.48550/ARXIV.2502.13923
- [2] Qiguang Chen, Libo Qin, Jinhao Liu, Dengyun Peng, Jiannan Guan, Peng Wang, Mengkang Hu, Yuhang Zhou, Te Gao, and Wanxiang Che. 2025. Towards Reasoning Era: A Survey of Long Chain-of-Thought for Reasoning Large Language Models. *CoRR* abs/2503.09567 (2025). arXiv:2503.09567 doi:10.48550/ARXIV.2503.09567
- [3] Yifei Chen, Guanting Dong, and Zhicheng Dou. 2025. Toward Effective Tool-Integrated Reasoning via Self-Evolved Preference Learning. arXiv:2509.23285 [cs.AI] <https://arxiv.org/abs/2509.23285>
- [4] DeepSeek-AI, Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, Xiao Bi, Xiaokang Zhang, Xingkai Yu, Yu Wu, Z. F. Wu, Zhibin Gou, Zhihong Shao, Zhuoshu Li, Ziyi Gao, Aixin Liu, Bing Xue, Bingxuan Wang, Bochao Wu, Bei Feng, Chengda Lu, Chenggang Zhao, Chengqi Deng, Chenyu Zhang, Chong Ruan, Damai Dai, Deli Chen, Dongjie Ji, Erhang Li, Fangyun Lin, Fucong Dai, Fuli Luo, Guangbo Hao, Guanting Chen, Guowei Li, H. Zhang, Han Bao, Hanwei Xu, Haocheng Wang, Honghui Ding, Huajian Xin, Huazuo Gao, Hui Qu, Hui Li, Jianzhong Guo, Jiashi Li, Jiawei Wang, Jingchang Chen, Jingyang Yuan, Junjie Qiu, Junlong Li, J. L. Cai, Jiaqi Ni, Jian Liang, Jin Chen, Kai Dong, Kai Hu, Kaige Gao, Kang Guan, Kexin Huang, Kuai Yu, Lean Wang, Lecong Zhang, Liang Zhao, Litong Wang, Liyue Zhang, Lei Xu, Leyi Xia, Mingchuan Zhang, Minghua Zhang, Minghui Tang, Meng Li, Miaojun Wang, Mingming Li, Ning Tian, Panpan Huang, Peng Zhang, Qiancheng Wang, Qinyu Chen, Qiusi Du, Ruiqi Ge, Ruisong Zhang, Ruizhe Pan, Runji Wang, R. J. Chen, R. L. Jin, Ruyi Chen, Shanghao Lu, Shangyan Zhou, Shanhuang Chen, Shengfeng Ye, Shiyu Wang, Shuiping Yu, Shunfeng Zhou, Shutong Pan, and S. S. Li. 2025. DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning. *CoRR* abs/2501.12948 (2025). arXiv:2501.12948 doi:10.48550/ARXIV.2501.12948
- [5] Guanting Dong, Licheng Bao, Zhongyuan Wang, Kangzhi Zhao, Xiaoxi Li, Jiajie Jin, Jinghan Yang, Hangyu Mao, Fuzheng Zhang, Kun Gai, Guorui Zhou, Yutao Zhu, Ji-Rong Wen, and Zhicheng Dou. 2025. Agentic Entropy-Balanced Policy Optimization. arXiv:2510.14545 [cs.LG] <https://arxiv.org/abs/2510.14545>
- [6] Guanting Dong, Yifei Chen, Xiaoxi Li, Jiajie Jin, Hongjin Qian, Yutao Zhu, Hangyu Mao, Guorui Zhou, Zhicheng Dou, and Ji-Rong Wen. 2025. Tool-Star: Empowering LLM-Brained Multi-Tool Reasoner via Reinforcement Learning. *CoRR* abs/2505.16410 (2025). arXiv:2505.16410 doi:10.48550/ARXIV.2505.16410
- [7] Guanting Dong, Jiajie Jin, Xiaoxi Li, Yutao Zhu, Zhicheng Dou, and Ji-Rong Wen. 2025. RAG-Critic: Leveraging Automated Critic-Guided Agentic Workflow for Retrieval Augmented Generation. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, ACL 2025, Vienna, Austria, July 27 - August 1, 2025, Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar (Eds.). Association for Computational Linguistics, 3551–3578. <https://aclanthology.org/2025.acl-long.179/>
- [8] Guanting Dong, Hangyu Mao, Kai Ma, Licheng Bao, Yifei Chen, Zhongyuan Wang, Zhongxia Chen, Jiazhen Du, Huiyang Wang, Fuzheng Zhang, Guorui Zhou, Yutao Zhu, Ji-Rong Wen, and Zhicheng Dou. 2025. Agentic Reinforced Policy Optimization. *CoRR* abs/2507.19849 (2025). arXiv:2507.19849 doi:10.48550/ARXIV.2507.19849
- [9] Runnan Fang, Shihao Cai, Baixuan Li, Jialong Wu, Guangyu Li, Wenbiao Yin, Xinyu Wang, Xiaobin Wang, Liangcai Su, Zhen Zhang, Shibin Yu, Zhengwei Tao, Yong Jiang, Pengjun Xie, Fei Huang, and Jingren Zhou. 2025. Towards General Agentic Intelligence via Environment Scaling. arXiv:2509.13311 [cs.CL] <https://arxiv.org/abs/2509.13311>
- [10] Jiazhan Feng, Shijue Huang, Xingwei Qu, Ge Zhang, Yujia Qin, Baoquan Zhong, Chengquan Jiang, Jinxin Chi, and Wanjuan Zhong. 2025. ReTool: Reinforcement Learning for Strategic Tool Use in LLMs. arXiv:2504.11536 [cs.CL] <https://arxiv.org/abs/2504.11536>
- [11] Zhiwei He, Tian Liang, Jiahao Xu, Qiuzhi Liu, Xingyu Chen, Yue Wang, Linfeng Song, Dian Yu, Zhenwen Liang, Wenxuan Wang, Zhuosheng Zhang, Rui Wang, Zhaopeng Tu, Haitao Mi, and Dong Yu. 2025. DeepMath-103K: A Large-Scale, Challenging, Decontaminated, and Verifiable Mathematical Dataset for Advancing Reasoning. (2025). arXiv:2504.11456 [cs.CL] <https://arxiv.org/abs/2504.11456>
- [12] Xinyi Hou, Yanjie Zhao, Shenao Wang, and Haoyu Wang. 2025. Model Context Protocol (MCP): Landscape, Security Threats, and Future Research Directions. *CoRR* abs/2503.23278 (2025). arXiv:2503.23278 doi:10.48550/ARXIV.2503.23278
- [13] Mengkang Hu, Tianxing Chen, Qiguang Chen, Yao Mu, Wenqi Shao, and Ping Luo. 2024. HiAgent: Hierarchical Working Memory Management for Solving Long-Horizon Agent Tasks with Large Language Model. arXiv:2408.09559 [cs.CL] <https://arxiv.org/abs/2408.09559>
- [14] Mengkang Hu, Bowei Xia, Yuran Wu, Ailing Yu, Yude Zou, Qiguang Chen, Shijian Wang, Jiarui Jin, Kexin Li, Wenxiang Jiao, Yuan Lu, and Ping Luo. 2025. Agent2World: Learning to Generate Symbolic World Models via Adaptive Multi-Agent Feedback. arXiv:2512.22336 [cs.AI] <https://arxiv.org/abs/2512.22336>
- [15] Mengkang Hu, Pu Zhao, Can Xu, Qingfeng Sun, Janguang Lou, Qingwei Lin, Ping Luo, and Saravan Rajmohan. 2025. AgentGen: Enhancing Planning Abilities for Large Language Model based Agent via Environment and Task Generation. arXiv:2408.00764 [cs.CL] <https://arxiv.org/abs/2408.00764>
- [16] Aaron Jaech, Adam Kalai, Adam Lerer, Adam Richardson, Ahmed El-Kishky, Aiden Low, Alec Helyar, Aleksander Madry, Alex Beutel, Alex Carney, et al. 2024. OpenAI o1 System Card. *arXiv preprint arXiv:2412.16720* (2024).
- [17] Dongfu Jiang, Yi Lu, Zhuofeng Li, Zhiheng Lyu, Ping Nie, Haozhe Wang, Alex Su, Hui Chen, Kai Zou, Chao Du, Tianyu Pang, and Wenhui Chen. 2025. VeriTool: Towards Holistic Agentic Reinforcement Learning with Tool Use. arXiv:2509.01055 [cs.AI] <https://arxiv.org/abs/2509.01055>
- [18] Bowen Jin, Hansi Zeng, Zhenrui Yue, Dong Wang, Hamed Zamani, and Jiawei Han. 2025. Search-R1: Training LLMs to Reason and Leverage Search Engines with Reinforcement Learning. *CoRR* abs/2503.09516 (2025). arXiv:2503.09516 doi:10.48550/ARXIV.2503.09516
- [19] Haolin Jin, Linghan Huang, Haipeng Cai, Jun Yan, Bo Li, and Huaming Chen. 2024. From LLMs to LLM-based Agents for Software Engineering: A Survey of Current, Challenges and Future. *CoRR* abs/2408.02479 (2024). arXiv:2408.02479 doi:10.48550/ARXIV.2408.02479
- [20] Jiajie Jin, Xiaoxi Li, Guanting Dong, Yuyao Zhang, Yutao Zhu, Yongkang Wu, Zhonghua Li, Ye Qi, and Zhicheng Dou. 2025. Hierarchical Document Refinement for Long-context Retrieval-augmented Generation. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, ACL 2025, Vienna, Austria, July 27 - August 1, 2025, Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar (Eds.). Association for Computational Linguistics, 3502–3520. <https://aclanthology.org/2025.acl-long.176/>
- [21] Jiajie Jin, Xiaoxi Li, Guanting Dong, Yuyao Zhang, Yutao Zhu, Zhao Yang, Hongjin Qian, and Zhicheng Dou. 2025. Decoupled Planning and Execution: A Hierarchical Reasoning Framework for Deep Search. *CoRR* abs/2507.02652 (2025). arXiv:2507.02652 doi:10.48550/ARXIV.2507.02652
- [22] Jiajie Jin, Yuyao Zhang, Yimeng Xu, Hongjin Qian, Yutao Zhu, and Zhicheng Dou. 2025. FinSight: Towards Real-World Financial Deep Research. arXiv:2510.16844 [cs.CL] <https://arxiv.org/abs/2510.16844>
- [23] Minki Kang, Wei-Ning Chen, Dongge Han, Huseyin A. Inan, Lukas Wutschitz, Yanzhi Chen, Robert Sim, and Saravan Rajmohan. 2025. ACON: Optimizing Context Compression for Long-horizon LLM Agents. arXiv:2510.00615 [cs.AI] <https://arxiv.org/abs/2510.00615>
- [24] Minghao Li, Yingxiu Zhao, Bowen Yu, Feifan Song, Hangyu Li, Haiyang Yu, Zhoujun Li, Fei Huang, and Yongbin Li. 2023. API-Bank: A Comprehensive Benchmark for Tool-Augmented LLMs. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing, EMNLP 2023, Singapore, December 6-10, 2023*, Houda Bouamor, Juan Pino, and Kalika Bali (Eds.). Association for Computational Linguistics, 3102–3116. doi:10.18653/V1/2023.EMNLP-MAIN.187
- [25] Xiaoxi Li, Guanting Dong, Jiajie Jin, Yuyao Zhang, Yujia Zhou, Yutao Zhu, Peitian Zhang, and Zhicheng Dou. 2025. Search-o1: Agentic Search-Enhanced Large Reasoning Models. *CoRR* abs/2501.05366 (2025). arXiv:2501.05366 doi:10.48550/ARXIV.2501.05366
- [26] Xiaoxi Li, Jiajie Jin, Guanting Dong, Hongjin Qian, Yutao Zhu, Yongkang Wu, Ji-Rong Wen, and Zhicheng Dou. 2025. WebThinker: Empowering Large Reasoning Models with Deep Research Capability. *CoRR* abs/2504.21776 (2025). arXiv:2504.21776 doi:10.48550/ARXIV.2504.21776
- [27] Xiaoxi Li, Jiajie Jin, Yujia Zhou, Yongkang Wu, Zhonghua Li, Ye Qi, and Zhicheng Dou. 2025. RetroLLM: Empowering Large Language Models to Retrieve Fine-grained Evidence within Generation. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, ACL 2025, Vienna, Austria, July 27 - August 1, 2025, Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar (Eds.). Association for Computational Linguistics, 16754–16779. <https://aclanthology.org/2025.acl-long.819/>
- [28] Xiaoxi Li, Yujia Zhou, and Zhicheng Dou. 2024. UniGen: A Unified Generative Framework for Retrieval and Question Answering with Large Language Models. In *Thirty-Eighth AAAI Conference on Artificial Intelligence, AAAI 2024, Thirty-Sixth Conference on Innovative Applications of Artificial Intelligence, IAAI 2024, Fourteenth Symposium on Educational Advances in Artificial Intelligence, EAAI 2024, February 20-27, 2024, Vancouver, Canada*, Michael J. Wooldridge, Jennifer G. Dy, and Sriraam Natarajan (Eds.). AAAI Press, 8688–8696. doi:10.1609/AAAI.V38I8.28714
- [29] Xuefeng Li, Haoyang Zou, and Pengfei Liu. 2025. ToRL: Scaling Tool-Integrated RL. arXiv:2503.23383 [cs.CL] <https://arxiv.org/abs/2503.23383>
- [30] Zhong-Zhi Li, Duzhen Zhang, Ming-Liang Zhang, Jiaxin Zhang, Zengyan Liu, Yuxuan Yao, Haotian Xu, Junhao Zheng, Pei-Jie Wang, Xiuyi Chen, Yingying Zhang, Fei Yin, Jiahua Dong, Zhijiang Guo, Le Song, and Cheng-Lin Liu. 2025. From System 1 to System 2: A Survey of Reasoning Large Language Models. *CoRR* abs/2502.17419 (2025). arXiv:2502.17419 doi:10.48550/ARXIV.2502.17419
- [31] Junteng Liu, Yunji Li, Chi Zhang, Jingyang Li, Aili Chen, Ke Ji, Weiye Cheng, Zijia Wu, Chengyu Du, Qidi Xu, Jiayuan Song, Zhengmao Zhu, Wenhui Chen, Pengyu Zhao, and Junxian He. 2025. WebExplorer: Explore and Evolve for Training Long-Horizon Web Agents. *CoRR* abs/2509.06501 (2025). arXiv:2509.06501 doi:10.

- 48550/ARXIV.2509.06501
- [32] Zichen Liu, Anya Sims, Keyu Duan, Changyu Chen, Simon Yu, Xiangxin Zhou, Haotian Xu, Shaopan Xiong, Bo Liu, Chenmian Tan, Chuen Yang Beh, Weixun Wang, Hao Zhu, Weiyang Shi, Diyi Yang, Michael Shieh, Yee Whye Teh, Wee Sun Lee, and Min Lin. 2025. GEM: A Gym for Agentic LLMs. arXiv:2510.01051 [cs.LG] <https://arxiv.org/abs/2510.01051>
 - [33] Grégoire Mialon, Clémentine Fourrier, Thomas Wolf, Yann LeCun, and Thomas Scialom. 2024. GAlA: a benchmark for General AI Assistants. In *The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7–11, 2024*. OpenReview.net. <https://openreview.net/forum?id=fibxvavhs3>
 - [34] OpenAI. 2025. Introducing deep research. <https://openai.com/index/introducing-deep-research>.
 - [35] Long Phan, Alice Gatti, Ziwen Han, Nathaniel Li, Josephina Hu, Hugh Zhang, Sean Shi, Michael Choi, Anish Agrawal, Arnav Chopra, Adam Khoja, Ryan Kim, Jason Hausenloy, Oliver Zhang, Mantas Mazeika, Daron Anderson, Tung Nguyen, Mobeen Mahmood, Fiona Feng, Steven Y. Feng, Haoran Zhao, Michael Yu, Varun Gangal, Chelsea Zou, Zihan Wang, Jessica P. Wang, Pawan Kumar, Oleksandr Pokutnyi, Robert Gerbic, Serguei Popov, John-Clark Levin, Mstyslav Kazakov, Johannes Schmitt, Geoff Galgon, Alvaro Sanchez, Yongki Lee, Will Yeadon, Scott Sauters, Marc Roth, Chidozie Agu, Søren Riis, Fabian Giska, Saiteja Utpala, Zachary Giboney, Gashaw M. Goshu, Joan of Arc Xavier, Sarah-Jane Crowson, Mohinder Maheshbhai Naiya, Noah Burns, Lennart Finke, Zerui Cheng, Hyunwoo Park, Francesco Fournier-Facio, John Wydallis, Mark Nandor, Ankit Singh, Tim Gehrunger, Jiaqi Cai, Ben McCarty, Darling Duclosel, Jungbae Nam, Jennifer Zampese, Ryan G. Hoerr, Aras Bacho, Gautier Abou Loume, Abdallah Galal, Hangrui Cao, Alexis C. Garretson, Damien Sileo, Qiuyu Ren, Doru Cojoc, Pavel Arkhipov, Usman Qazi, Lianghui Li, Sumeet Motwani, Christian Schröder de Witt, Edwin Taylor, Johannes Veith, Eric Singer, Taylor D. Hartman, Paolo Rissone, Jaehyeok Jin, Jack Wei Lun Shi, Chris G. Willcocks, Joshua Robinson, Aleksandar Mikov, Ameya Prabhu, Longke Tang, Xavier Alapont, Justine Leon Uro, Kevin Zhou, Emily de Oliveira Santos, Andrey Pupasov Maksimov, Edward Vendrow, Kengo Zenitani, Julien Guillo, Yuqi Li, Joshua Vendrow, Vladyslav Kuchkin, and Ng Ze-An. 2025. Humanity's Last Exam. CoRR abs/2501.14249 (2025). arXiv:2501.14249 doi:10.48550/ARXIV.2501.14249
 - [36] Yiwei Qin, Xuefeng Li, Haoyang Zou, Yixiu Liu, Shijie Xia, Zhen Huang, Yixin Ye, Weizhe Yuan, Hector Liu, Yuanzhi Li, et al. 2024. O1 Replication Journey: A Strategic Progress Report—Part 1. *arXiv preprint arXiv:2410.18982* (2024).
 - [37] Yujia Qin, Shihao Liang, Yining Ye, Kunlun Zhu, Lan Yan, Yaxi Lu, Yankai Lin, Xin Cong, Xiangru Tang, Bill Qian, Sihao Zhao, Lauren Hong, Runchu Tian, Ruobing Xie, Jie Zhou, Mark Gerstein, Dahai Li, Zhiyuan Liu, and Maosong Sun. 2024. ToolLLM: Facilitating Large Language Models to Master 16000+ Real-world APIs. In *The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7–11, 2024*. OpenReview.net. <https://openreview.net/forum?id=dHng2O0Jjr>
 - [38] Changle Qu, Sunhao Dai, Xiaochi Wei, Hengyi Cai, Shuaiqiang Wang, Dawei Yin, Jun Xu, and Ji-Rong Wen. 2025. From Exploration to Mastery: Enabling LLMs to Master Tools via Self-Driven Interactions. In *The Thirteenth International Conference on Learning Representations, ICLR 2025, Singapore, April 24–28, 2025*. OpenReview.net. <https://openreview.net/forum?id=QKBu1BOAwd>
 - [39] Changle Qu, Sunhao Dai, Xiaochi Wei, Hengyi Cai, Shuaiqiang Wang, Dawei Yin, Jun Xu, and Ji-Rong Wen. 2025. Tool learning with large language models: a survey. *Frontiers Comput. Sci.* 19, 8 (2025), 198343. doi:10.1007/S11704-024-40678-2
 - [40] Qwen, ., An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, Huan Lin, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Yang, Jiayi Yang, Jingren Zhou, Junyang Lin, Kai Dang, Keming Lu, Keqin Bao, Kexin Yang, Le Yu, Mei Li, Mingfeng Xue, Pei Zhang, Qin Zhu, Rui Men, Runji Lin, Tianhao Li, Tingyu Xia, Xingzhang Ren, Xuancheng Ren, Yang Fan, Yang Su, Yichang Zhang, Yu Wan, Yuqiong Liu, Zeyu Cui, Zhenru Zhang, and Zihan Qiu. 2024. Qwen2.5 Technical Report. arXiv:2412.15115 [cs.CL] <https://arxiv.org/abs/2412.15115>
 - [41] Guangming Sheng, Chi Zhang, Zilingfeng Ye, Xibin Wu, Wang Zhang, Ru Zhang, Yanghua Peng, Haibin Lin, and Chuan Wu. 2024. HybridFlow: A Flexible and Efficient RLHF Framework. *arXiv preprint arXiv: 2409.19256* (2024).
 - [42] Zhengliang Shi, Yuhang Wang, Lingyong Yan, Pengjie Ren, Shuaiqiang Wang, Dawei Yin, and Zhaochun Ren. 2025. Retrieval Models Aren't Tool-Savvy: Benchmarking Tool Retrieval for Large Language Models. In *Findings of the Association for Computational Linguistics, ACL 2025, Vienna, Austria, July 27–August 1, 2025*, Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar (Eds.). Association for Computational Linguistics, 24497–24524. <https://aclanthology.org/2025.findings-acl.1258/>
 - [43] Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. 2023. Reflexion: language agents with verbal reinforcement learning. In *Advances in Neural Information Processing Systems 36: Annual Conference on Neural Information Processing Systems 2023, NeurIPS 2023, New Orleans, LA, USA, December 10–16, 2023*, Alice Oh, Tristan Naumann, Amir Globerson, Kate Saenko, Moritz Hardt, and Sergey Levine (Eds.). http://papers.nips.cc/paper_files/paper/2023/hash/1b44b878bb782e6954cd88628510e90-Abstract-Conference.html
 - [44] Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. 2024. Reflexion: Language agents with verbal reinforcement learning. *Advances in Neural Information Processing Systems* 36 (2024).
 - [45] Mohit Shridhar, Xingdi Yuan, Marc-Alexandre Côté, Yonatan Bisk, Adam Trischler, and Matthew J. Hausknecht. 2021. ALFWorld: Aligning Text and Embodied Environments for Interactive Learning. In *9th International Conference on Learning Representations, ICLR 2021, Virtual Event, Austria, May 3–7, 2021*. OpenReview.net. <https://openreview.net/forum?id=0lOX0YcCdTn>
 - [46] Xiaoshuai Song, Haoqi Chang, Guanting Dong, Yutao Zhu, Zhicheng Dou, and Ji-Rong Wen. 2026. EnvScaler: Scaling Tool-Interactive Environments for LLM Agent via Programmatic Synthesis. arXiv:2601.05808 [cs.CL] <https://arxiv.org/abs/2601.05808>
 - [47] Yifan Song, Weimin Xiong, Dawei Zhu, Cheng Li, Ke Wang, Ye Tian, and Sujian Li. 2023. RestGPT: Connecting Large Language Models with Real-World Applications via RESTful APIs. CoRR abs/2306.06624 (2023). arXiv:2306.06624 doi:10.48550/ARXIV.2306.06624
 - [48] Weiwei Sun, Miao Lu, Zhan Ling, Kang Liu, Xuesong Yao, Yiming Yang, and Jiecao Chen. 2025. Scaling Long-Horizon LLM Agent via Context-Folding. arXiv:2510.11967 [cs.CL] <https://arxiv.org/abs/2510.11967>
 - [49] Jiejun Tan, Zhicheng Dou, Yan Yu, Jiehan Cheng, Qiang Ju, Jian Xie, and Ji-Rong Wen. 2025. HierSearch: A Hierarchical Enterprise Deep Search Framework Integrating Local and Web Searches. arXiv:2508.08088 [cs.IR] <https://arxiv.org/abs/2508.08088>
 - [50] Zhengwei Tao, Jialong Wu, Wenbiao Yin, Junkai Zhang, Baixuan Li, Haiyang Shen, Kuan Li, Liwen Zhang, Xinyu Wang, Yong Jiang, Pengjun Xie, Fei Huang, and Jingren Zhou. 2025. WebShaper: Agentially Data Synthesizing via Information-Seeking Formalization. CoRR abs/2507.15061 (2025). arXiv:2507.15061 doi:10.48550/ARXIV.2507.15061
 - [51] Qwen Team. 2024. Qwq: Reflect deeply on the boundaries of the unknown. *Hugging Face* (2024).
 - [52] Lei Wang, Wanyu Xu, Yihuai Lan, Zhiqiang Hu, Yunshi Lan, Roy Ka-Wei Lee, and Ee-Peng Lim. 2023. Plan-and-Solve Prompting: Improving Zero-Shot Chain-of-Thought Reasoning by Large Language Models. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2023, Toronto, Canada, July 9–14, 2023*, Anna Rogers, Jordan L. Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, 2609–2634. doi:10.18653/V1/2023.ACL-LONG.147
 - [53] Renxi Wang, Xudong Han, Lei Ji, Shu Wang, Timothy Baldwin, and Haonan Li. 2025. ToolGen: Unified Tool Retrieval and Calling via Generation. In *The Thirteenth International Conference on Learning Representations, ICLR 2025, Singapore, April 24–28, 2025*. OpenReview.net. <https://openreview.net/forum?id=XLMMmowdY>
 - [54] Shijian Wang, Runhao Fu, Siyi Zhao, Qingqin Zhan, Xingjian Wang, Jiarui Jin, Yuan Lu, Hanqian Wu, and Cunjian Chen. 2025. Synthetic Curriculum Reinforces Compositional Text-to-Image Generation. *arXiv preprint arXiv:2511.18378* (2025).
 - [55] Shijian Wang, Jiarui Jin, Xingjian Wang, Linxin Song, Runhao Fu, Hecheng Wang, Zongyuan Ge, Yuan Lu, and Xuelian Cheng. 2025. Video-Thinker: Sparking "Thinking with Videos" via Reinforcement Learning. *arXiv preprint arXiv:2510.23473* (2025).
 - [56] Xingyao Wang, Yangyi Chen, Lifan Yuan, Yizhe Zhang, Yunzhu Li, Hao Peng, and Heng Ji. 2024. Executable Code Actions Elicit Better LLM Agents. In *Forty-first International Conference on Machine Learning, ICML 2024, Vienna, Austria, July 21–27, 2024*. OpenReview.net. <https://openreview.net/forum?id=jJ9BoXAFfa>
 - [57] Yinuo Wang, Mining Tan, Wenxiang Jiao, Xiaoxi Li, Hao Wang, Xuanyu Zhang, Yuan Lu, and Weiming Dong. 2026. TourPlanner: A Competitive Consensus Framework with Constraint-Gated Reinforcement Learning for Travel Planning. arXiv:2601.04698 [cs.AI] <https://arxiv.org/abs/2601.04698>
 - [58] Yinuo Wang, Likun Wang, Yuxuan Jiang, Wenjun Zou, Tong Liu, Xujie Song, Wenxuan Wang, Liming Xiao, Jiang Wu, Jingliang Duan, and Shengbo Eben Li. 2024. Diffusion Actor-Critic with Entropy Regulator. arXiv:2405.15177 [cs.LG] <https://arxiv.org/abs/2405.15177>
 - [59] Zihan Wang, Kangrui Wang, Qinqing Wang, Pingyue Zhang, Linjie Li, Zhengyuan Yang, Xing Jin, Kefan Yu, Minh Nhat Nguyen, Licheng Liu, Eli Gottlieb, Yiping Lu, Kyunghyun Cho, Jiajun Wu, Li Fei-Fei, Lijuan Wang, Yejin Choi, and Manling Li. 2025. RAGEN: Understanding Self-Evolution in LLM Agents via Multi-Turn Reinforcement Learning. CoRR abs/2504.20073 (2025). arXiv:2504.20073 doi:10.48550/ARXIV.2504.20073
 - [60] Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed H. Chi, Quoc V. Le, and Denny Zhou. 2022. Chain-of-Thought Prompting Elicits Reasoning in Large Language Models. In *Advances in Neural Information Processing Systems 35: Annual Conference on Neural Information Processing Systems 2022, NeurIPS 2022, New Orleans, LA, USA, November 28–December 9, 2022*, Sanmi Koyejo, S. Mohamed, A. Agarwal, Danielle Belgrave, K. Cho, and A. Oh (Eds.). http://papers.nips.cc/paper_files/paper/2022/hash/9d5609613524ecf4f15af0f7b31abca4-Abstract-Conference.html
 - [61] Jialong Wu, Baixuan Li, Runnan Fang, Wenbiao Yin, Liwen Zhang, Zhengwei Tao, Dingchu Zhang, Zekun Xi, Yong Jiang, Pengjun Xie, Fei Huang, and Jingren Zhou. 2025. WebDancer: Towards Autonomous Information Seeking Agency.

- CoRR abs/2505.22648 (2025). arXiv:2505.22648 doi:10.48550/ARXIV.2505.22648
- [62] Shitao Xiao, Zheng Liu, Peitian Zhang, Niklas Muennighoff, Defu Lian, and Jian-Yun Nie. 2024. C-Pack: Packed Resources For General Chinese Embeddings. In *Proceedings of the 47th International ACM SIGIR Conference on Research and Development in Information Retrieval, SIGIR 2024, Washington DC, USA, July 14-18, 2024*, Grace Hui Yang, Hongning Wang, Sam Han, Claudia Hauff, Guido Zuccon, and Yi Zhang (Eds.). ACM, 641–649. doi:10.1145/3626772.3657878
- [63] Yang Xiao, Mohan Jiang, Jie Sun, Keyu Li, Jifan Lin, Yumin Zhuang, Ji Zeng, Shijie Xia, Qishuo Hua, Xuefeng Li, Xiaojie Cai, Tongyu Wang, Yue Zhang, Liming Liu, Xia Wu, Jinlong Hou, Yuan Cheng, Wenjie Li, Xiang Wang, Dequan Wang, and Pengfei Liu. 2025. LIMI: Less is More for Agency. arXiv:2509.17567 [cs.AI] <https://arxiv.org/abs/2509.17567>
- [64] Zhenghai Xue, Longtao Zheng, Qian Liu, Yingru Li, Xiaosen Zheng, Zejun Ma, and Bo An. 2025. Simpletr: End-to-end reinforcement learning for multi-turn tool-integrated reasoning. *arXiv preprint arXiv:2509.02479* (2025).
- [65] An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, Chujie Zheng, Dayiheng Liu, Fan Zhou, Fei Huang, Feng Hu, Hao Ge, Haoran Wei, Huan Lin, Jialong Tang, Jian Yang, Jianhong Tu, Jianwei Zhang, Jian Yang, Jiayi Yang, Jingren Zhou, Junyang Lin, Kai Dang, Keqin Bao, Kexin Yang, Le Yu, Lianghao Deng, Mei Li, Mingfeng Xue, Mingze Li, Pei Zhang, Peng Wang, Qin Zhu, Rui Men, Ruize Gao, Shixuan Liu, Shuang Luo, Tianhao Li, Tianyi Tang, Wenbiao Yin, Xingzhang Ren, Xinyu Wang, Xinyu Zhang, Xuancheng Ren, Yang Fan, Yang Su, Yichang Zhang, Yinger Zhang, Yu Wan, Yuyuqiong Liu, Zekun Wang, Zeyu Cui, Zhenru Zhang, Zhipeng Zhou, and Zihan Qiu. 2025. Qwen3 Technical Report. *CoRR abs/2505.09388* (2025). arXiv:2505.09388 doi:10.48550/ARXIV.2505.09388
- [66] Wenkai Yang, Shuming Ma, Yankai Lin, and Furu Wei. 2025. Towards Thinking-Optimal Scaling of Test-Time Compute for LLM Reasoning. *CoRR abs/2502.18080* (2025). arXiv:2502.18080 doi:10.48550/ARXIV.2502.18080
- [67] Shunyu Yao, Howard Chen, John Yang, and Karthik Narasimhan. 2022. WebShop: Towards Scalable Real-World Web Interaction with Grounded Language Agents. In *Advances in Neural Information Processing Systems 35: Annual Conference on Neural Information Processing Systems 2022, NeurIPS 2022, New Orleans, LA, USA, November 28 - December 9, 2022*, Sanmi Koyejo, S. Mohamed, A. Agarwal, Danielle Belgrave, K. Cho, and A. Oh (Eds.). http://papers.nips.cc/paper_files/paper/2022/hash/82ad13ec01f9fe44c01cb91814fd7b8c-Abstract-Conference.html
- [68] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. 2022. React: Synergizing reasoning and acting in language models. *arXiv preprint arXiv:2210.03629* (2022).
- [69] Junjie Ye, Zhengyin Du, Xuesong Yao, Weijian Lin, Yufei Xu, Zehui Chen, Zaiyuan Wang, Sining Zhu, Zhiheng Xi, Siyu Yuan, Tao Gui, Qi Zhang, Xuanjing Huang, and Jiecao Chen. 2025. ToolHop: A Query-Driven Benchmark for Evaluating Large Language Models in Multi-Hop Tool Use. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2025, Vienna, Austria, July 27 - August 1, 2025*, Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar (Eds.). Association for Computational Linguistics, 2995–3021. <https://aclanthology.org/2025.acl-long.150/>
- [70] Aohan Zeng, Mingdao Liu, Rui Lu, Bowen Wang, Xiao Liu, Yuxiao Dong, and Jie Tang. 2024. AgentTuning: Enabling Generalized Agent Abilities for LLMs. In *Findings of the Association for Computational Linguistics, ACL 2024, Bangkok, Thailand and virtual meeting, August 11-16, 2024*, Lun-Wei Ku, Andre Martins, and Vivek Srikumar (Eds.). Association for Computational Linguistics, 3053–3077. doi:10.18653/V1/2024.FINDINGS-ACL.181
- [71] Qianchi Zhang, Hainan Zhang, Liang Pang, Hongwei Zheng, Yongxin Tong, and Zhiming Zheng. 2025. Less is More: Compact Clue Selection for Efficient Retrieval-Augmented Generation Reasoning. arXiv:2502.11811 [cs.CL] <https://arxiv.org/abs/2502.11811>
- [72] Yuxiang Zheng, Dayuan Fu, Xiangkun Hu, Xiaojie Cai, Lyumanshan Ye, Pengrui Lu, and Pengfei Liu. 2025. DeepResearcher: Scaling Deep Research via Reinforcement Learning in Real-world Environments. *arXiv preprint arXiv:2504.03160* (2025).

Appendix

A Datasets

A.1 Training Data

We collected a diverse training dataset spanning four task categories to instill comprehensive agent capabilities.

- **General Tool-Use:** We sample 1k instances for labeled-tool scenarios and 1k for tool-retrieval from the ToolBench [37] training set. This data is intended to instill a generalized ability to use diverse tools and leverage large toolsets through retrieval.

- **Real-World Interaction:** We utilize 500 instances from ALF-World [45] and 500 from WebShop [67], sampled from their training sets, to teach the model to interact effectively with environments, manage state transitions, and achieve user goals.
- **Deep Research:** We include 200 instances from WebDancer [61] and 500 from WebShaperQA [50] to enhance the model’s proficiency in using web search and page browsing for in-depth information gathering.
- **Mathematical Reasoning:** We collect 0.9k problems from the DeepMath dataset [11] to strengthen the model’s ability to use code as a tool for complex mathematical computations.

A.2 Benchmarks

We conduct extensive experiments on a wide range of benchmarks, including general tool-use and downstream applications.

General Tool-Use. These benchmarks encompass a broad range of distinct tools (from tens to over ten thousand), thus offering a testbed for evaluating different approaches to toolset scaling.

- **ToolBench [37]:** A large-scale benchmark containing over 16,000 real-world REST APIs spanning 49 categories. Test subsets include 100 test cases, designed to evaluate LLMs in both single-tool and complex multi-tool scenarios.
- **API-Bank [24]:** A comprehensive benchmark for tool-augmented LLMs. It features a runnable evaluation system with 73 API tools and a large training set (over 2,200 dialogues across 2,211 APIs from 1,008 domains), assessing LLMs’ capabilities in planning, retrieving, and calling APIs.
- **TMDB [47]:** A sub-scenario of RestBench focused on the TMDB movie database, consisting of 100 questions that utilize 54 local tools and require an average of 2.3 sequential API calls.
- **Spotify [47]:** A sub-scenario of RestBench simulating a Spotify music player, featuring 57 questions and 40 local tools, demanding an average of 2.6 sequential API calls to complete the tasks.
- **ToolHop [69]:** A multi-hop reasoning dataset comprising 995 complex questions. It leverages 3,912 locally executable tools and requires between 3 to 7 sequential tool calls per task.

Downstream Applications. These benchmarks test the capability of different approaches in handling complex real-world tasks, which often require the use of domain-specific toolsets.

- **ALFWorld [45]:** A benchmark for simple Embodied AI tasks set in a text environment. Agents must complete objectives using a finite set of low-level embodied actions (eg., move, take) to test navigation and object manipulation.
- **WebShop [67]:** A challenging online shopping environment that provides 12,087 crowd-sourced tasks over a catalog of 1.18 million products. Agents interact with the simulated e-commerce website using core APIs: search[Query] and choose[Text Button].
- **GAIA [33]:** A complex benchmark for General AI Assistants, consisting of 466 real-world questions (with a 300-question held-out test set). It requires the flexible application of a broad general-purpose toolset including web browsing, code execution, multi-modal processing, and file handling.
- **Humanity’s Last Exam (HLE) [35]:** A benchmark featuring 2,500 highly difficult, multi-disciplinary questions (graduate-level). It primarily evaluates the model’s intrinsic deep reasoning and multi-modal understanding capabilities, as the questions are designed to be insoluble by simple external search tools.

Table 6: A case study of DeepAgent for multi-step, multi-tool tasks on ToolBench, showing only key actions.

Example of DeepAgent on a Multi-Tool Task
User Query: I'm organizing a film festival and I need assistance in finding the best films. Can you search for videos related to 'documentary' on Vimeo? Additionally, fetch the related people in the 'cinema' category to invite them as guest speakers. Finally, provide me with a streaming link for a YouTube video with the ID 'UxxajLWwzqY'.
DeepAgent Processing Output
Okay, I need to solve three tasks: find documentaries on Vimeo, find cinema-related people on Vimeo, and get a YouTube streaming link. I'll start by searching for the right tools. <code><tool_search></code> Vimeo video search for documentaries <code></tool_search></code> The search returned several tools. <code>search_videos</code> and <code>get_videos_by_tag</code> seem most relevant. I'll use the search tool first to find documentaries. <code><tool_call></code> {"name": "search_videos", "arguments": {"query": "documentary", "sort": "most_liked"}} <code></tool_call></code> The tool returned a list of popular documentary videos. Now for the second task: finding people in the 'cinema' category. Since there is no direct tool to search for people, I will try searching for videos with the 'cinema' tag and identify their creators. <code><tool_call></code> {"name": "get_videos_by_tag", "arguments": {"tag": "cinema", "sort": "most_liked"}} <code></tool_call></code> This call returned videos tagged with 'cinema' and included user information for each video, such as "John Doe" and "Jane Smith", who are potential guest speakers. Finally, for the YouTube streaming link, I'll first verify the video ID exists using the YouTube search tool. <code><tool_call></code> {"name": "search_for_youtube_v3", "arguments": {"part": "snippet", "type": "video", "q": "UxxajLWwzqY"}} <code></tool_call></code> The tool confirmed the video ID is valid. I can now construct the streaming link. I will now consolidate all the information to answer the user's query. Final Answer: 1. Documentaries on Vimeo: Found nature and space documentaries. 2. Cinema Guest Speakers: John Doe, Jane Smith with Vimeo profiles. 3. YouTube Link: https://www.youtube.com/watch?v=UxxajLWwzqY

B Implementation Details

For DeepAgent, we use QwQ-32B [51] as the main reasoning model for the results in Table 1 and Table 2, and Qwen3-30B-A3B-Thinking-2507 [65] with Qwen3-235B-A22B-Thinking-2507 [65] in Table 5. We use Qwen2.5-32B-Instruct [40] as the auxiliary model for (1) filtering lengthy tool search results and execution outputs (this is also applied to all baselines), (2) simulating RapidAPIs during ToolPO training, and (3) generating folded memory from interaction history. For the baselines, we use either QwQ-32B or Qwen2.5-32B-Instruct as the backbone model. Text generation for all models uses a maximum of 81,920 tokens, with a temperature of 0.7, top_p of 0.8, top_k of 20, and a repetition penalty of 1.05. The maximum number of actions is set to 50.

Web search and page browsing are implemented using the Google Serper API and Jina Reader API, respectively. The VQA tool is based on Qwen2.5-VL-32B-Instruct [1], which takes a question and an image as input and outputs a model-generated response. Tool retrieval is performed using bge-large-en-v1.5 [62]. All tool documentation follows the standard OpenAI function definition format: {"name": "...", "description": "...", "parameters": {"type": "object", "properties":

{"param1": {"type": "...", "description": "...", "required": [{"param1"}]}}. This format is used for building the toolset index and for all prompts given to the agents.

Training consists of 100 steps of ToolPO with a batch size of 64, $\lambda_1 = \lambda_2 = 1$, rollout size $K = 8$, and a maximum sequence length of 32,768. The maximum number of actions is 50. The training framework is based on VeRL [41] for multi-node distributed training. All experiments are conducted on 64 NVIDIA H20-141GB GPUs.

C Memory Schema

Our brain-inspired memory architecture contains three components: episodic, working, and tool memory. To support stable memory folding and reduce information loss, we define each component with a fixed JSON schema, enabling reliable parsing and use of compressed memories for long-horizon reasoning.

Episodic Memory Schema. Episodic memory records high-level task progression (milestones, decisions, outcomes) to preserve long-term context. The format is: {"task_description": "A general summary of what the reasoning history has been doing and the overall goals it has been striving for.", "key_events": [{"step": "step number", "description": "A detailed description of the specific action taken, decision made, or milestone achieved at this step, including relevant context and reasoning behind the choice.", "outcome": "A detailed account of the direct result, observation, or feedback received from this action or decision, including any new information gained or changes in the task state."}], "current_progress": "A general summary of the current progress of the task, including what has been completed and what is left to be done."}

Working Memory Schema. Working memory captures the immediate goal, active challenges, and next actions to maintain continuity across folds. The format is: {"immediate_goal": "A clear summary of the current subgoal—what you are actively working toward at this moment.", "current_challenges": "A concise summary of the main obstacles or difficulties you are presently encountering.", "next_actions": [{"type": "tool_call or planning ordecision", "description": "Anticipate and describe the next concrete action you intend to take to advance the task."}]}

Tool Memory Schema. Tool memory consolidates tool-use experience (success rates, effective parameters, common errors) and derives reusable rules. The format is: {"tools_used": [{"tool_name": "string", "success_rate": "float", "effective_parameters": [{"param1", "param2"}], "common_errors": [{"error_type1", "error_type2"}], "response_pattern": "description of typical output", "experience": "Reflect and summarize your experience, including both successes and failures."}], "derived_rules": [{"When X condition occurs, prefer tool Y", "Tool Z works best with parameter A set to B"}]}

D Case Study

To illustrate the effectiveness of our DeepAgent framework in handling complex, multi-step tasks that require coordinated use of multiple tools, we present a detailed case in Table 6. This example demonstrates how DeepAgent autonomously navigates tool selection, executes sequential actions, and synthesizes results to provide comprehensive solutions to user queries.